

Learning Nonlinearities


Introduction

Consider the XOR operator. In plain language, $y = \text{XOR}(x_1, x_2)$ lights up when exactly one of x_1 or x_2 is true, but not both. That is, if we have $x_1 \in \{0, 1\}$ and $x_2 \in \{0, 1\}$, $y = \text{XOR}(x_1, x_2)$ takes its value according to:

x_1	x_2	y
0	0	0
0	1	1
1	0	1
1	1	0

The Classification Problem

Since $y \in \{0, 1\}$, we can think of this setup in terms of a classification problem with the predictors x_1 and x_2 .




The big problem:

We cannot come up with a decision boundary that is linear in x_1 and x_2 which separates the classes!!

Feature Engineering

One idea is to try feature engineering. That is, construct $f_1 = \phi_1(x_1, x_2)$ and $f_2 = \phi_2(x_1, x_2)$ such that the classes are linearly separable in the space spanned by (f_1, f_2) . We want to end up with something like:



But how should we pick $\phi_1(\cdot, \cdot)$ and $\phi_2(\cdot, \cdot)$?

Since this example is really simple, we can just think a bit and arrive through trial and error at some good choices.

Example


For example, take:

$$f_1 = \phi_1(x_1, x_2) = \max(x_1 + x_2, 0)$$
$$f_2 = \phi_2(x_1, x_2) = \max(x_1 + x_2 - 1, 0)$$

Then we have:

x_1	x_2	f_1	f_2	y
0	0	0	0	0
0	1	1	0	1
1	0	1	0	1
1	1	2	1	0


Visually, now we have:



Now they are linearly separable!! :)

Connection to Neural Networks/Representation Learning

This example of feature engineering also illustrates how neural networks work!
Here is the idea:



If we take $W = \begin{bmatrix} 0 & 1 & 1 \\ -1 & 1 & 1 \end{bmatrix}$ and the ReLU activation function,

$$\sigma(x) = \text{ReLU}(x) = x \mathbf{1}_{\{x>0\}}(x) = \max(x, 0)$$

then the constructed features are:

$$f_1 = \sigma(w_{10} + w_{11}x_1 + w_{12}x_2) = (x_1 + x_2)_+ = \phi_1(x_1, x_2)$$

$$f_2 = \sigma(w_{20} + w_{21}x_1 + w_{22}x_2) = (x_1 + x_2 - 1)_+ = \phi_2(x_1, x_2)$$

The last step is just to train a logistic regression on them and learn the $\hat{\beta}$ vector.
Knowing the $\hat{\beta}$ coefficients is equivalent to knowing the decision boundary!

A Couple Things to Note

1. This problem was simple enough so that we could just guess the values of W . In general, we will have to learn them!
2. Sometimes classes are not linearly separable because of noise, sometimes because of a true structural nonlinearity (or both). In this case, we observe the values of $y = \text{XOR}(x_1, x_2)$ without noise, so we know it is a true structural nonlinearity. For structural nonlinearities, neural networks can help a lot; they can't do much about the noise problem, though.
3. The concept that the neural network is trying to learn how to represent the features f is one of the big ideas of neural networks. Sometimes you'll hear this referred to as "representation learning."