

AdaBoost

Recall that, in general, boosting is an additive stagewise model, i.e.,

$$f(x) = \sum_{m=1}^M \alpha_m T_m(x).$$

Where $T_m(x) = \{-1, +1\}$ and the true labels $y \in \{-1, +1\}$.

Our prediction \hat{y} is defined as:

$$\hat{y} = \text{sign}(f(x)).$$

It's useful for us to define the *margin* as:

$$m(x) = y \cdot f(x).$$

Claim: An obs (x, y) is classified correctly iff the margin is positive, $m(x) > 0$.

Proof: If $m(x) = y \cdot f(x) > 0$ then y and $f(x)$ have the same sign. Then if $m(x) > 0 \implies \hat{y} = \text{sign}(f(x)) \implies \hat{y} = y$.

If $y \cdot f(x) < 0 \implies \hat{y}$ is the opposite sign of y . $\implies \hat{y} \neq y$.

Things to notice:

1. $f(x) = 0$ is the decision boundary.
2. The margin can be thought of as like a residual for binary classification.
3. Loss functions for binary classification problem can be written in terms of the margin.

Claim: Log loss for binary classification (starting from the $y^* = \{0, 1\}$ logistic regression setup) can be written as:

$$L(y, f(x)) = \log[1 + e^{-y f(x)}],$$

if we switch to the convention of taking $y \in \{-1, +1\}$ and the classification rule is $\hat{y} = \text{sign}(f(x))$.

Proof: Consider the $y^* = \{0, 1\}$ logistic regression scenario. Recall $f(x)$ is just the linear predictor in this case: $f(x) = ax + b$, and thus the soft prediction is,

$$\hat{y}^* = \text{sigmoid}(f(x)) = \frac{1}{1 + e^{-f(x)}}.$$

The binary log-loss in general takes the form:

$$L(y^*, \hat{y}^*) = -[y^* \log(\hat{y}^*) + (1 - y^*) \log(1 - \hat{y}^*)].$$

Here's the main idea we use to simplify this expression:

$$y^* = 1 \implies y = 1, \text{ and } y^* = 0 \implies y = -1.$$

So in each case,

$$L(y^*, \hat{y}^*) = \begin{cases} -\log(\hat{y}^*), & \text{if } y^* = 1 \iff y = 1, \\ -\log(1 - \hat{y}^*), & \text{if } y^* = 0 \iff y = -1. \end{cases}$$

If $y = 1$,

$$-\log(\hat{y}^*) = -\log[(1 + e^{-f(x)})^{-1}] = \log[1 + e^{-f(x)}].$$

And if $y = -1$,

$$\begin{aligned} -\log(1 - \hat{y}^*) &= -\log\left(1 - \frac{1}{1 + e^{-f(x)}}\right) = -\log\left(1 - \frac{e^{f(x)}}{e^{f(x)} + 1}\right) = \\ -\log\left(\frac{e^{f(x)} + 1 - e^{f(x)}}{e^{f(x)} + 1}\right) &= -\log\left(\frac{1}{1 + e^{f(x)}}\right) = \log[1 + e^{f(x)}] = \log[1 + e^{-y f(x)}]. \end{aligned}$$

Examples of Losses for Binary Classification

Log-loss \Rightarrow


$$\ell(y, f(x)) = \log(1 + e^{-y f(x)}).$$

Exp. \Rightarrow

$$L(y, f(x)) = e^{-y f(x)}.$$

Hinge \Rightarrow

$$H^H(y, f(x)) = \max(0, 1 - y \cdot f(x)).$$



Notice:

- Loss functions penalize negative margins. How do each of these penalize? Look at the asymptotic behavior for negative margins:

Exp. $\rightarrow e^{-m(x)}$ (exponential penalty),

LogLoss $\rightarrow \log(1 + e^{-m(x)}) \approx$ linear for large negative margins,

Hinge \rightarrow linear for negative margins.

What is AdaBoost doing?

AdaBoost is:

1. optimizing exponential loss, and
2. fitting additive models in steps.

Consider the additive model:

$$f(x) = \sum_{m=1}^M \beta_m T_m(x).$$

Problem: Find β_1, \dots, β_m and T_1, \dots, T_m such that

$$\frac{1}{n} \sum_{i=1}^n L(y_i, f(x_i)) = \frac{1}{n} \sum_{i=1}^n L\left(y_i, \sum_{m=1}^n \beta_m T_m(x)\right)$$

is minimized. Finding all trees at once is intractable, so we work in stages.

Stage 1: Find β and T such that

$$\frac{1}{n} \sum_{i=1}^n L(y_i, \beta T(x_i)) = \frac{1}{n} \sum_{i=1}^n e^{-y_i \beta T(x_i)} \quad \text{is minimized.}$$

First,

$$\begin{aligned} \frac{1}{n} \sum_{i=1}^n L(y_i, \beta T(x_i)) &= \frac{1}{n} \sum_{i=1}^n e^{-y_i \beta T(x_i)} \\ &= \frac{1}{n} \left(\sum_{\text{margin}=1} e^{-\beta} + \sum_{\text{margin}=-1} e^{\beta} \right). \end{aligned}$$

Define

$$\text{err} = \frac{\#\{\text{incorrectly classified points}\}}{n}.$$

and

$$Q(\beta) = \frac{1}{n} \sum_{i=1}^n L(y_i, \beta T(x_i)) = \text{err} \cdot e^{\beta} + (1 - \text{err}) \cdot e^{-\beta}.$$

Minimize with respect to β ,

$$\frac{dQ}{d\beta} = \text{err} e^{\beta} - (1 - \text{err}) e^{-\beta} = 0.$$

$$\begin{aligned} \text{err } e^{2\beta} &= 1 - \text{err.} \implies e^{2\beta} = \frac{1 - \text{err}}{\text{err}}. \\ 2\beta &= \log\left(\frac{1 - \text{err}}{\text{err}}\right) \implies \beta^* = \beta_1 = \frac{1}{2} \log\left(\frac{1 - \text{err}}{\text{err}}\right). \end{aligned}$$

Here β_1 is the “amount of say” that the first tree has!

Stage m for $m > 2$:

At the $(m-1)^{th}$ stage, we have the current function:

$$f_{m-1}(x) = \sum_{j=1}^{m-1} \beta_j T_j(x).$$

We want to find T_m such that

$$\frac{1}{n} \sum_{i=1}^n L(y_i | f_{m-1}(x_i) + \beta_m T_m(x_i)) = \frac{1}{n} \sum_{i=1}^n e^{-y_i [f_{m-1}(x_i) + \beta_m T_m(x_i)]}$$

is minimized.

We can rewrite:

$$\frac{1}{n} \sum_{i=1}^n e^{-y_i [f_{m-1}(x_i) + \beta_m T_m(x_i)]} = \frac{1}{n} \sum_{i=1}^n e^{-y_i f_{m-1}(x_i)} e^{-y_i \beta_m T_m(x_i)} = \frac{1}{n} \sum_{i=1}^n w_i e^{-y_i \beta_m T_m(x_i)}$$

where

$$w_i = e^{-y_i f_{m-1}(x_i)}.$$

Expanding this expression:

$$\begin{aligned} \frac{1}{n} \sum_{i=1}^n w_i e^{-y_i \beta_m T_m(x_i)} &= \frac{1}{n} \left[\sum_{i: y_i = T_m(x_i)} w_i e^{-\beta_m} + \sum_{i: y_i \neq T_m(x_i)} w_i e^{\beta_m} \right] \\ &= \frac{1}{n} \left[(e^\beta - e^{-\beta}) \sum_{i=1}^n w_i \mathbf{1}[y_i \neq T_m(x_i)] + e^{-\beta} \sum_{i=1}^n w_i \right] \end{aligned}$$

To minimize this, we want

$$\sum_{i=1}^n w_i \mathbf{1}[y_i \neq T_m(x_i)]$$

to be as small as possible. That is:

$$T_m^* = \arg \min_T \sum_{i=1}^n w_i \mathbf{1}[y_i \neq T(x_i)].$$

This is the tree that minimizes the weighted error.

The best β_m is, as before,

$$\beta_m = \frac{1}{2} \log\left(\frac{1 - \text{err}_m}{\text{err}_m}\right), \quad \text{where } \text{err}_m = \frac{\sum_{i=1}^n w_i \mathbf{1}[y_i \neq T_m(x_i)]}{\sum_{i=1}^n w_i}.$$

When we update our weights for the next tree, we do:

$$w_i \leftarrow w_i \exp[-\beta_m y_i T_m(x_i)] = \begin{cases} w_i e^{-\beta_m}, & \text{if } y_i = T_m(x_i), \\ w_i e^{\beta_m}, & \text{if } y_i \neq T_m(x_i). \end{cases}$$

Showing equivalence to the α_m formulation:

We can redefine some quantities to show that this is totally equivalent to the α_m formulation. Define:

$$\alpha_m = 2\beta_m,$$

so that

$$\alpha_m = \log\left(\frac{1 - \text{err}_m}{\text{err}_m}\right),$$

and redefine

$$w_i \leftarrow w_i \exp[\alpha_m \mathbf{1}(y_i \neq T_m(x_i))].$$

Then we're done.

Big Takeaway: Finally our fitted boosted tree looks like:

$$f(x) = \sum_{m=1}^M \beta_m T_m(x_i) = \frac{1}{2} \sum_{m=1}^M \alpha_m T_m(x_i) \implies \hat{y} = \text{sign}(f(x)).$$

and since at the end of the day we are only using the sign of $f(x)$ for the hard prediction, multiplying by 2 doesn't change the final prediction \hat{y} .