

# Advanced Machine Learning

Content-based Rec. Systems  
& AdaBoost

# Outline


- Go over quiz
- Content-based Rec. Systems
- Review of decision trees
- Boosting question
- Ensemble models
- Weighted decision stumps
- AdaBoost

# **Content-based recommendations**


# Classification of Recom. Systems

- **Content-based systems:**
  - analysis of the attributes of the items (movies, books) to generate predictions
  - build user profiles based on the **content features** of the items rated by the user
  - match user profile attributes with against the attributes of the content object
- **Collaborative filtering:**
  - recommend items based on similarity measures between users and/or items
  - items recommended to a user are those preferred by similar users

## COLLABORATIVE FILTERING



## CONTENT-BASED FILTERING



# What type of system?

You run a website that recommends blog post to users.

You don't collect user ids or cookies.


*You want to recommend a next article to read. What type of system would you recommend with this type of data?*

*Why?*

# Ingredients for Content-based Rec.

- Item representation
  - Based on content
- User representation
  - Based on items “liked” by user
- Similarity metric / Machine learning model
- Sometimes ratings

# Example: music recommendation



# Item profile for music recommendation

- Audio Content Analysis
  - Features can be extracted from any audio file
    - Beat, timbre, tonal features
  - Learn high level descriptors from low-level features via machine learning (e.g. deep learning)
    - Learning genre (house, bass, funk, DnB, jazz, classical, metal)
    - Learning semantic categories: not\_danceable, gender\_male, mood\_not\_happy
  - Many libraries will do some of this already

# Item profile for music recommendation

- Audio Content Analysis
- Text Content Analysis
  - Text processing of user generated context and lyrics

# Item profile for movie recommendation

- The set of actors of the movie.
- The director.
- The year in which the movie was made.
  - Some viewers prefer old movies, others watch only the latest releases.
- The genre or general type of movie.
  - Some viewers like only comedies, others dramas or romances.

# Item profile for movie recommendation

- The set of actors of the movie.
- The director.
- The year in which the movie was made.
  - Some viewers prefer old movies, others watch only the latest releases.
- The genre or general type of movie.
  - Some viewers like only comedies, others dramas or romances.

Can you think about other features? Can we make features from reviews? How do we get a user profile?

# Item profile for movie recommendation

- Source
  - Author, publisher
- Location
  - Movies (only interesting for a region), pictures can be tagged to location
  - Represented with latitude, longitude or country/state/city
- Image features
- Audio features
- Application specific

# User profile for movie recommendations

- **Content-based profile:** summary of the profile of the items these user liked / purchased

Example **Netflix recommendations**

**Item profile:**

$x_1$  = has Julia Roberts,  
 $x_2$  = directed by Lars von Trier,  
 $x_3$  = is horror,  
 $x_4$  = is a comedy

**User profile:**

$x_1$  = probability that the user likes movies with Julia Roberts,  
 $x_2$  = probability that the user likes movies directed by Lars von Trier,  
 $x_3$  = probability that the user that likes horror movies,  
 $x_4$  = probability that the user likes comedies

# User profile for movie recommendations

- Demographics
- Declared Interests
- User current location (from IP address)
- Usage-based features
  - Last time visit, frequency of visits (weekly, monthly), frequency per device
- Search history
- Item set
  - Set of items for which the user showed interest (e.g. clicked, shared, liked)

# Recommending Items to Users based on content: $kNN$

- Compute profile vectors for users and items
- Find a similarity measure and compute similarity between users and items
- Recommend items with *high similarity* to the user prof.
  - K-nearest neighbors

# Recommending Items to Users based on content: $kNN$

- Scaling up
  - When you have too many users and items it is not feasible to compute similarity between all of them
  - Locality-sensitive-hashing techniques can be used to place item profiles in buckets
  - Given a user, easy to find buckets with high similarity to the user

# Recommending Items to Users: *training a model*

- Compute profile vectors for users and items
- Train a model using the feature vector to predict observed ratings
  - Actual ratings, clicks, likes or a combination
  - Regression to predict numerical ratings
  - Classification to predict prob of click
  - This approach could include features that are not content-based

# Advantages of Content-based Recommendation

- You don't need data on other users
- You don't have **cold-start** problem for a new item
  - Able to recommend new and unpopular items
- You can provide explanations of recommended items by listing content features that caused an item to be recommended

# Challenges with Content-based Recommendation

- Constructing the feature vector could be a difficult task (need domain knowledge)
- New genres are hard to recognize
- Some kind of items are not amenable to easy feature extraction (movies, music)
- Hard to exploit quality of judgements of other users

# Helpful References

- <https://www.sciencedirect.com/science/article/pii/S1110866515000341>
- [https://research.googleblog.com/2018/01/the-google-brain-team-looking-back-on\\_12.html](https://research.googleblog.com/2018/01/the-google-brain-team-looking-back-on_12.html)
- <https://nlp.stanford.edu/manning/talks/Simons-Institute-Manning-2017.pdf>
- [http://www.cp.jku.at/tutorials/mrs\\_recsys\\_2017/slides.pdf](http://www.cp.jku.at/tutorials/mrs_recsys_2017/slides.pdf)

# Boosting


Required Reading: Ch. 18 of  
Probabilistic Machine Learning  
by Kevin Patrick Murphy

Emphasis on Sect. 18.5 - Boosting

# Review of Decision Trees


# Review of Decision Trees

- Popular machine learning algorithm
- Building block for random forests and gradient boosting
- Interpretable
- “Bushy trees” tend to
  - Overfit
  - Have high variance (Sensitive to small changes in the training data)
  - Low bias (Fit well the training data)





# Review of Decision Trees

- Bushy trees tend to
  - Overfit
  - Have high variance (Sensitive to small changes in the training data)
  - Have low bias
- Shallow trees tend to
  - Underfit
  - Have high bias



Decision tree problem  $T(x) = (1)1_{(x \in R_1)} + (-1)1_{(x \in R_2)}$   
 $+ (-1)1_{(x \in R_3)} + (1)1_{(x \in R_4)}$




Problem: Find the best split that separates the circles from triangles. Splits must be parallel to axis.

# Decision tree problem

Problem: Find the best split that separates the circles from triangles. Splits must be parallel to axis.

This is a really hard problem!!




# Impurity

How to formalize the decision tree problem?

Find the splits that minimize impurity.

# Measures of impurity

For a two class problem:



# Decision tree pseudocode

1. Place the best attribute ( $x_j$  and value  $a$ ) of the dataset at the root of the tree
2. Split the training set into subsets:  
 $\{x_j \leq a\}$  and  $\{x_j > a\}$
3. Repeat step 1 and step 2 on each subset until you find leaf nodes or a stopping condition is met

# Binary decision stumps

Let's  $x = (x_1, \dots, x_D)$  where each feature  $x_j$  is a continuous variable.

$$T(x; \theta) = \begin{cases} 1, & \text{if } x_j > \alpha \\ -1, & \text{otherwise} \end{cases}$$

Or

$$T(x; \theta) = \begin{cases} 1, & \text{if } x_j \leq \alpha \\ -1, & \text{otherwise} \end{cases}$$



$$\begin{matrix} +1 & & -1 \\ \hline \underline{\underline{=}} & & \underline{\underline{=}} \end{matrix}$$



# Tree notation

Trees partition the space into *disjoint* regions  $R_j, j \dots J$ . For  $x$  in particular region  $R_j$  the value of  $T(x)$  constant. A tree can be formally expressed as:

$$\mathcal{T}(x; \theta) = \sum_{j=1}^J \beta_j \mathbf{1}_{[x \in R_j]}$$

where  $\theta$  is the set of parameters for the trees  $\theta = \{\{R_j\}_1^J; \{\beta_j\}_1^J\}$ .

# Tree notation: Exercise



$$R_1 = \{x < 0.43\} \cap \{y < 0.47\}$$

$$T(x; \theta) = \sum_{j=1}^J \beta_j \mathbf{1}_{[x \in R_j]}$$

What are the parameters for this tree?


# The Boosting Question

# Binary classification with hard predictions


Given training observations,  $(x^{(1)}, y^{(1)}), \dots, (x^{(N)}, y^{(N)})$

Let  $\hat{y}^{(i)}$  be either **1** or **-1**.

# Weak classifiers



Logistic regression w. simple features




Shallow decision trees




Decision stumps

- Fast to fit
- Does not overfit (low variance)
- Tends to underfit: does not fit our data well (high bias)

# Finding the right classifier: Learning curve



# How to get a strong classifier?



- Add polynomial features or more features
- Ensemble
  - Random Forest
- Is there another way??


# The Boosting Question

“Can a set of weak learners be combined to create a strong learner?” (Kearns and Valiant, 1988)

Weak learner: a learner with  $> 50\%$  accuracy

F class err.

# Ensemble: each classifier “votes” on prediction



$x = (Income = 120K, \text{Credit} = \text{Bad}, \text{Savings} = 50K, \text{Market} = \text{Good})$

$$G(x) = \text{sign}(\alpha_1 G_1(x) + \alpha_2 G_2(x) + \alpha_3 G_3(x) + \alpha_4 G_4(x))$$

Handwritten annotations with blue ink:

- Four blue underlines are placed under the  $\alpha$  coefficients:  $\underline{\alpha}_1$ ,  $\underline{\alpha}_2$ ,  $\underline{\alpha}_3$ , and  $\underline{\alpha}_4$ .
- A blue curly brace is drawn under the four terms  $\underline{\alpha}_1 G_1(x) + \underline{\alpha}_2 G_2(x) + \underline{\alpha}_3 G_3(x) + \underline{\alpha}_4 G_4(x)$ .

# Ensemble classifier

- Base learner that outputs 1 or -1
- Ensemble model
  - M classifiers
  - M coefficients
- Prediction

$$\hat{y} = \text{sign} \left( \sum_{m=1}^M \alpha_m \underline{G}_m(x) \right)$$

# Prediction with ensemble:

## What is the prediction in this example?



$$G_1(x) = +1$$

$$G_2(x) = -1$$

$$G_3(x) = -1$$

$$G_4(x) = +1$$

$x = (Income = \$120K, Credit = Bad, Savings = \$50K, Market = Good)$

$$G(x) = \text{sign}(\alpha_1 G_1(x) + \alpha_2 G_2(x) + \alpha_3 G_3(x) + \alpha_4 G_4(x))$$

where  $\alpha_1 = 2, \alpha_2 = 1.5, \alpha_3 = 1.5, \alpha_4 = 0.5$

$$G(x) = \text{sign}(2 - 1.5 - 1.5 + 0.5) = \text{sign}(-0.5) = -1$$

# Preview of AdaBoost

---



## Algorithm 1 Adaboost

---



```
1: procedure ADABOOST
2:   Maintain a weight of each point in a training set  $\{(x^{(i)}, y^{(i)})\}_{i=1}^N$ 
3:   Initialize weights  $w_i = \frac{1}{N}, i \in [1 \cdots N]$ 
4:   for iter = 1 to  $M$  do
5:     Apply base learner to weighted examples.
6:     Increase weights of misclassified examples.
7:   end for
8:   Combine models by weighted voting.
9: end procedure
```

---



# AdaBoost: first decision stump




# AdaBoost: learned weights



# AdaBoost: Second decision stump



# AdaBoost Ensemble



# Learning a decision stump

Classification error = 
$$\frac{\text{\# incorrect predictions}}{\text{\# examples}}$$



Find the **feature** and the **threshold** that will minimize the classification error.

How would you generalize classification error to weighted data?

# Learning weighted decision stumps

Weighted  
cross  
err.

$$\theta^* = \arg \min_{\theta} \frac{\sum_{i=1}^N w_i \mathbb{1}_{[y^{(i)} \neq T(x^{(i)}; \theta)]}}{\sum_{i=1}^N w_i}$$

A fancy notation for **weighted classification error**

Exercise


$$\frac{2/5}{6/5} = \frac{1}{3}$$

Consider the following weighted dataset

$$x^{(1)} = (0, 0), y^{(1)} = 1,$$

$$x^{(2)} = (1, 1), y^{(2)} = -1,$$

$$x^{(3)} = (2, 2), y^{(3)} = 1,$$




Find a stump that minimizes the weighted classification error.

# Exercise

$$\frac{1/5}{6/5} = 1/6$$

1. Consider the following dataset.

$x = (x_1, x_2)$	$y$	weight
(0, 0)	-1	$\frac{2}{10}$
(0, 1)	-1	$\frac{2}{10}$
(1, 1)	-1	$\frac{3}{10}$
(1, 3)	1	$\frac{1}{10}$
(2, 1)	1	$\frac{2}{10}$



- (a) Compute the weighted classification error of the following classifier. Which points are misclassified?

$$g(x) = \begin{cases} 1, & \text{if } x_2 \geq 0.5 \\ -1, & \text{otherwise} \end{cases}$$

- (b) Find a stump that minimizes the weighted classification error. What is the best weighted classification error?

---

**Algorithm 2** Adaboost

---

1: **procedure** ADABOOT

2:    Initialize weights  $w_i = \frac{1}{N}, i \in [1 \dots N]$

3:    **for**  $m = 1$  to  $M$  **do**

4:     Fit a tree classifier  $T_m(x) = T(x; \theta_m)$  to training data using  
weights  $w_i$

5:     Compute

$$err_m = \frac{\sum_{i=1}^N w_i \mathbb{1}_{[y^{(i)} \neq T_m(x^{(i)})]}}{\sum_{i=1}^N w_i}$$

6:     Compute  $\alpha_m = \log(\frac{1-err_m}{err_m})$

7:     Set  $w_i = \underbrace{w_i \cdot \exp}_{\leftarrow} \underbrace{[\alpha_m \cdot \mathbb{1}_{[y^{(i)} \neq T_m(x^{(i)})]}]}_{\leftarrow}, i = 1, \dots, N$

8:    **end for**

9:     $f(x) = \sum_{m=1}^M \alpha_m T_m(x)$

10:   Return  $F(x) = \underbrace{\text{sign}(f(x))}_{\leftarrow}$

11: **end procedure**

---

# Computing classifier weights

$$\alpha_m = \log\left(\frac{1 - \text{err}_m}{\text{err}_m}\right)$$

error	alpha
0.01	$\log\left(\frac{99}{.01}\right) \approx \log(99) \approx$
0.2	$\log\left(\frac{0.8}{0.2}\right) \approx \log(4) \approx$
0.5	$\log\left(\frac{0.5}{0.5}\right) = \log(1) = 0$

We want the classifier to have large weight if error is small and a small weight if error is not so good.

# Reweighting data to focus on mistakes

Increase the weights  
on incorrectly classified  
observations

$$w_i = w_i \cdot \exp[\alpha_m \cdot 1_{[y^{(i)} \neq T_m(x^{(i)})]}]$$

$$w_i = \begin{cases} w_i & \text{if } y^{(i)} = T_m(x^{(i)}) \\ w_i \cdot \exp[\alpha_m] & \text{otherwise} \end{cases}$$

	$T_m(x) = y$	err of $T_m$	New weight
obs 1	Correct	<del>0.1</del> 0.2	$w_1$
2	Correct	<del>0.5</del> 0.2	$w_2$
3	Incorrect	0.2	$4w_3$

## Exercise:

Run AdaBoost  
by hand on this  
dataset for 2  
iterations

$x_1$ <b>temperature</b>	$x_2$ <b>wind</b>	$y$ <b>play tennis</b>
hot	F	no
hot	T	no
hot	F	yes
mild	F	yes
cool	F	yes
cool	T	no

# References

- Probabilistic Machine Learning, Ch. 18
- <https://www.coursera.org/learn/ml-classification/lecture/3ywWA/the-boosting-question>
- The Elements of Statistical Learning
- <https://www.coursera.org/learn/competitive-data-science/lecture/br2ze/practical-guide>
- [https://www.utdallas.edu/~nrr150130/cs7301/2016fa/lects/Lecture\\_10\\_Ensemble.pdf](https://www.utdallas.edu/~nrr150130/cs7301/2016fa/lects/Lecture_10_Ensemble.pdf).